




PBFT

Byzantine Fault Tolerance Consensus Algorithm Simulation



Our team:

Cai Huanghao

Zeng Xianrong

Wang Ying


Yu Xiaochen



Distributing data across multiple nodes requires a reliable consensus algorithm.

The Problem: Malicious nodes, message delays, and network failures can all prevent agreement.

Solution: PBFT (Practical Byzantine Fault Tolerance) offers a solution to these challenges.





GOAL OF PROJECT

Construct a PBFT simulator to visualize and understand the consensus process.

The simulator will allow users to customize parameters, configure network conditions, and observe how PBFT works.

Visualize the network topology, monitor node interactions, and analyze consensus performance.



Technology stack

Python 3.6+
with tkinter, random,
collections.defaultdict



ABOUT PBFT

The Byzantine Generals Problem:

Imagine a group of generals who must agree on a battle plan.

Some generals (nodes) may be faulty or lie, making agreement impossible.

A Fault-Tolerant Consensus: PBFT

solves this challenge by ensuring that honest generals can still reach an agreement.

Core Threshold

safety threshold

$$f \leq \lfloor \frac{N-1}{3} \rfloor$$

the maximum number of malicious nodes the system can tolerate while maintaining consistency

consensus threshold


$$T = \lfloor \frac{2N}{3} \rfloor + 1$$

the minimum number of identical valid messages required to transition between protocol phases




The Execution Flow


- **Request:** Client sends a request to the Primary node.
- **Pre-Prepare:** Primary sends a pre-prepare message to all Backup nodes.
- **Prepare:** Backup nodes send prepare messages to each other.
- **Commit:** Backup nodes send commit messages to each other.
- **Reply:** Primary and Backup nodes send reply messages to the client.





Realization process





Node Class Attributes:



Attribute	Description
node_id	Unique node identifier
is_malicious	Malicious node flag (True=malicious, False=honest)
preprepare_val	Stores received pre-prepare message (correct value)
prepare_cnt	Counts received prepare messages (dict: proposal=count)
commit_cnt	Counts received commit messages (dict: proposal=count)
prepared	Whether prepared state is reached (True if prepare threshold met)
committed	Whether committed state is reached (True if commit threshold met)
result	Final consensus result (None if uncommitted)

Node Class Attributes:

```
# Node class: Simulate a single node in PBFT system
class Node:
    def __init__(self, node_id, is_malicious=False):
        self.node_id = node_id
        self.is_malicious = is_malicious # Whether the node is malicious
        self.preprepare_val = None # Pre-prepare message (always correct value)
        self.prepare_cnt = defaultdict(int) # Count of prepare messages received
        self.commit_cnt = defaultdict(int) # Count of commit messages received
        self.prepared = False # Whether reach prepared state
        self.committed = False # Whether reach committed state
        self.result = None # Final consensus result
```


PBFT_Simulator Class Attributes:



Attribute	Description
N	Total number of consensus nodes
f	Number of malicious nodes in the network
honest_num	Count of honest nodes (total - malicious)
safe	PBFT security check: $f \leq (N-1)//3$ (True=valid)
threshold	Consensus approval threshold ($\geq 2/3$ nodes)
malicious_ids	Randomly selected malicious node IDs (1~N-1, node 0 is honest primary)
nodes	List storing all Node class instances
honest_node_ids	List of all honest node identifiers
value	Correct proposed value for consensus (block hash)

PBFT_Simulator Class Attributes:

```
# PBFT Simulator class: Control the entire consensus process
class PBFT_Simulator:
    def __init__(self, total_nodes, malicious_num):
        self.N = total_nodes          # Total number of nodes
        self.f = malicious_num        # Number of malicious nodes
        self.honest_num = total_nodes - malicious_num # Number of honest nodes

        # Standard PBFT security condition:  $f \leq (N-1)/3$ 
        self.safe = (self.f <= (self.N - 1) // 3)

        # Consensus threshold: more than 2/3 of total nodes (rounded up)
        self.threshold = (2 * self.N) // 3 + 1

        # Create nodes: Node 0 is primary node (always honest)
        # Randomly select f nodes from 1~N-1 as malicious nodes
        malicious_ids = random.sample([i for i in range(1, self.N)], self.f) if self.N > 1 else []
        self.nodes = []
        self.honest_node_ids = [] # Store all honest node IDs
        for i in range(self.N):
            is_mal = (i in malicious_ids)
            self.nodes.append(Node(i, is_malicious=is_mal))
            if not is_mal:
                self.honest_node_ids.append(i)

        self.value = "block_hash_123456" # Correct proposal value
```

PBFT_Simulator Class Methods:

Method	Description
pre_prepare()	Pre-prepare phase: All nodes receive correct proposal value
prepare()	Prepare phase: Honest nodes broadcast real msg; malicious broadcast fake msg; count & verify prepare state
commit()	Commit phase: Prepared honest nodes broadcast commit msg; count & verify commit state
run()	Execute full PBFT consensus: run pre-prepare → prepare → commit → show result
show_result()	Display final consensus result of honest nodes; judge consensus success

```
# Run the complete PBFT consensus process
def run(self):
    print("==== PBFT Consensus Simulation Started =====")
    print("Total nodes: (self.N) | Malicious nodes: (self.M) | Honest nodes: (self.honest_num)")
    print("PBFT Security Threshold: T = (N+1)/3 = (self.T) + (self.M-1)/3 = (self.safeT)")
    print("Consensus Threshold: (self.T) of total nodes: (self.threshold)")

    self.pre_prepare()
    self.prepare()
    self.commit()
    self.show_result()
```

```
# 1. Pre-prepare phase: All nodes (including malicious) receive correct value
def pre_prepare(self):
    print("==== Pre-prepare Phase (All nodes receive correct proposal) =====")
    for node in self.nodes:
        node.preprepare_val = self.value # Key rule: All nodes get correct value!
        status = 'Malicious' if node.is_malicious else 'Honest'
        print(f"Node {node.node_id} [{status}] received: {node.preprepare_val}")
```

```
# 2. Prepare phase: Malicious nodes broadcast fake messages
def prepare(self):
    print("==== Prepare Phase (Malicious nodes broadcast fake messages) =====")
    # Step 1: All nodes broadcast messages
    for sender in self.nodes:
        # Determine broadcast content
        if sender.is_malicious:
            # Malicious nodes broadcast fake message
            send_msg = 'fake' + (self.value, sender.node_id)
        else:
            # Honest nodes broadcast real message
            send_msg = self.value

    # Broadcast to all nodes
    for receiver in self.nodes:
        if receiver.is_malicious:
            continue # Malicious nodes don't count messages
        receiver.prepare_cnt[send_msg] += 1 # Receiver counts messages
```

```
# Step 2: Check if each honest node reaches committed state
print(f"Prepare Phase Statistics (Honest Nodes Only):")
print("Total honest nodes: (self.honest_num)")
print("Honest node ID list: (self.honest_node_ids)")
print("Prepared state check for each honest node:")
for node in self.honest_nodes:
    if node.is_malicious:
        continue
    # Count correct messages
    correct_cnt = node.prepare_cnt.get(self.value, 0)
    # Check if each threshold
    if correct_cnt == self.threshold:
        node.prepared = True
        print(f"Honest Node {node.node_id}: Correct messages {correct_cnt} > threshold {self.threshold}. Prepared state achieved!")
    else:
        node.prepared = False
        print(f"Honest Node {node.node_id}: Correct messages {correct_cnt} < threshold {self.threshold}. Prepared state not achieved!")
```

```
# 3. Commit phase: Only prepared honest nodes broadcast commit messages
def commit(self):
    print("==== Commit Phase =====")
    # Step 1: Broadcast commit messages
    for sender in self.nodes:
        if sender.is_malicious or not sender.prepared:
            continue # Malicious/prepared nodes don't broadcast
        # Broadcast correct proposal to all honest nodes
        for receiver in self.honest_nodes:
            if receiver.is_malicious:
                continue
            receiver.commit_cnt[self.value] += 1
```

```
# Step 2: Check if each honest node reaches committed state
print("Commit Phase Statistics (Honest Nodes Only):")
for node in self.honest_nodes:
    if node.is_malicious:
        continue
    commit_cnt = node.commit_cnt.get(self.value, 0)
    print(f"Node {node.node_id}: Received commit messages {commit_cnt}")
    if commit_cnt == self.threshold:
        node.committed = True
        node.result = self.value
    else:
        node.committed = False
        node.result = None
```

```
# Show final consensus result
def show_result(self):
    print("==== Final Consensus Result =====")
    honest_results = []
    for node in self.honest_nodes:
        if node.is_malicious:
            continue
        print(f"Honest Node {node.node_id} Final Result: {node.result}")
        honest_results.append(node.result)

    # Judge if consensus succeeded
    if len(set(honest_results)) == 1 and honest_results[0] is not None:
        print("✅ Consensus Success! All honest nodes reached agreement")
    else:
        print("❌ Consensus Failed! Honest nodes did not reach agreement")
```


PBFT_Simulator Class Methods:

Method

Description

pre_prepare()

Pre-prepare phase: All nodes receive correct proposal value

```
# 1. Pre-prepare phase: All nodes (including malicious) receive correct value
def pre_prepare(self):
    print("\n==== Pre-prepare Phase (All nodes receive correct proposal) =====")
    for node in self.nodes:
        node.preprepare_val = self.value # Key rule: All nodes get correct value!
        status = "Malicious" if node.is_malicious else "Honest"
        print(f"Node {node.node_id} [{status}] received: {node.preprepare_val}")
```

PBFT_Simulator Class Methods:

Method	Description
pre_prepare()	Pre-prepare phase: All nodes receive correct proposal value
prepare()	Prepare phase: Honest nodes broadcast real msg; malicious broadcast fake msg; count & verify prepare state
commit()	Commit phase: Prepared honest nodes broadcast commit msg; count & verify commit state
run()	Execute full PBFT consensus: run pre-prepare → prepare → commit → show result
show_result()	Display final consensus result of honest nodes; judge consensus success

```
# Run the complete PBFT consensus process
def run(self):
    print("==== PBFT Consensus Simulation Started =====")
    print("Total nodes: (self.N) | Malicious nodes: (self.M) | Honest nodes: (self.honest_num)")
    print("PBFT Security Threshold: T = (N+1)/3 = (self.T) + (self.S) + (self.A)")
    print("Consensus Threshold: (self.T) of total nodes: (self.threshold)")

    self.pre_prepare()
    self.prepare()
    self.commit()
    self.show_result()
```

```
# 1. Pre-prepare phase: All nodes (including malicious) receive correct value
def pre_prepare(self):
    print("==== Pre-prepare Phase (All nodes receive correct proposal) =====")
    for node in self.nodes:
        node.preprepare_val = self.value # Key rule: All nodes get correct value!
        status = 'Malicious' if node.is_malicious else 'Honest'
        print(f"Node {node.node_id} [{status}] received: {node.preprepare_val}")
```

```
# 2. Prepare phase: Malicious nodes broadcast fake messages
def prepare(self):
    print("==== Prepare Phase (Malicious nodes broadcast fake messages) =====")
    # Step 1: All nodes broadcast messages
    for sender in self.nodes:
        # Determine broadcast content
        if sender.is_malicious:
            # Malicious nodes broadcast fake message
            send_msg = 'fake' + (self.value, sender.node_id)
        else:
            # Honest nodes broadcast real message
            send_msg = self.value

    # Broadcast to all nodes
    for receiver in self.nodes:
        if receiver.is_malicious:
            continue # Malicious nodes don't count messages
        receiver.prepare_cnt[send_msg] += 1 # Receiver counts messages
```

```
# Step 2: Check if each honest node reaches committed state
print(f"Prepare Phase Statistics (Honest Nodes Only):")
print("Total honest nodes: (self.honest_num)")
print("Honest node ID list: (self.honest_node_ids)")
print("Prepared state check for each honest node:")
for node in self.honest_nodes:
    if node.is_malicious:
        continue
    # Count correct messages
    correct_cnt = node.prepare_cnt.get(self.value, 0)
    # Check if each threshold
    if correct_cnt == self.threshold:
        node.prepared = True
        print(f"Honest Node {node.node_id}: Correct messages {correct_cnt} > threshold {self.threshold}. Prepared state achieved!")
    else:
        node.prepared = False
        print(f"Honest Node {node.node_id}: Correct messages {correct_cnt} < threshold {self.threshold}. Prepared state not achieved!")
```

```
# 3. Commit phase: Only prepared honest nodes broadcast commit messages
def commit(self):
    print("==== Commit Phase =====")
    # Step 1: Broadcast commit messages
    for sender in self.nodes:
        if sender.is_malicious or not sender.prepared:
            continue # Malicious/prepared nodes don't broadcast
        # Broadcast correct proposal to all honest nodes
        for receiver in self.honest_nodes:
            if receiver.is_malicious:
                continue
            receiver.commit_cnt[self.value] += 1
```

```
# Step 2: Check if each honest node reaches committed state
print("Commit Phase Statistics (Honest Nodes Only):")
for node in self.honest_nodes:
    if node.is_malicious:
        continue
    commit_cnt = node.commit_cnt.get(self.value, 0)
    print(f"Node {node.node_id}: Received commit messages {commit_cnt}")
    if commit_cnt == self.threshold:
        node.committed = True
        node.result = self.value
    else:
        node.committed = False
        node.result = None
```

```
# Show final consensus result
def show_result(self):
    print("==== Final Consensus Result =====")
    honest_results = []
    for node in self.honest_nodes:
        if node.is_malicious:
            continue
        print(f"Honest Node {node.node_id} Final Result: {node.result}")
        honest_results.append(node.result)

    # Judge if consensus succeeded
    if len(set(honest_results)) == 1 and honest_results[0] is not None:
        print("Consensus Success! All honest nodes reached agreement")
    else:
        print("Consensus Failed! Honest nodes did not reach agreement")
```

PBFT_Simulator Class Methods:

Method

Description

prepare()

Prepare phase: Honest nodes broadcast real msg; malicious broadcast fake msg; count & verify prepare state

```
# 2. Prepare phase: Malicious nodes broadcast fake messages
def prepare(self):
    print("\n==== Prepare Phase (Malicious nodes broadcast fake messages) =====")
    # Step 1: All nodes broadcast messages
    for sender in self.nodes:
        # Determine broadcast content
        if sender.is_malicious:
            # Malicious node: broadcast fake message
            send_msg = f"fake_{self.value}_{sender.node_id}"
        else:
            # Honest node: broadcast real message
            send_msg = self.value

    # Broadcast to all nodes
    for receiver in self.nodes:
        if receiver.is_malicious:
            continue # Malicious nodes don't count messages
        receiver.prepare_cnt[send_msg] += 1 # Receiver counts messages
```

```
# Step 2: Show honest node info and check Prepared state
print("\n[Prepare Phase Statistics (Honest Nodes Only):")
print(f"Total honest nodes: {self.honest_num}")
print(f"Honest node ID list: {self.honest_node_ids}")
print("\nPrepared state check for each honest node:")
for node in self.nodes:
    if node.is_malicious:
        continue
    # Count correct messages
    correct_cnt = node.prepare_cnt.get(self.value, 0)
    # Check if reach threshold
    if correct_cnt >= self.threshold:
        node.prepared = True
        print(f"→ Node {node.node_id}: Correct messages={correct_cnt} ≥ threshold({self.threshold}), Prepared state achieved")
    else:
        node.prepared = False
        print(f"→ Node {node.node_id}: Correct messages={correct_cnt} < threshold({self.threshold}), Prepared state not achieved")
```


PBFT_Simulator Class Methods:

Method	Description
pre_prepare()	Pre-prepare phase: All nodes receive correct proposal value
prepare()	Prepare phase: Honest nodes broadcast real msg; malicious broadcast fake msg; count & verify prepare state
commit()	Commit phase: Prepared honest nodes broadcast commit msg; count & verify commit state
run()	Execute full PBFT consensus: run pre-prepare → prepare → commit → show result
show_result()	Display final consensus result of honest nodes; judge consensus success

```
# Run the complete PBFT consensus process
def run(self):
    print("==== PBFT Consensus Simulation Started =====")
    print("Total nodes: (self.N) | Malicious nodes: (self.M) | Honest nodes: (self.honest_num)")
    print("PBFT Security Threshold: T = (N+1)/3 = (self.T) + (self.N-1)/3 = (self.safeT)")
    print("Consensus Threshold: (self.T) of total nodes: (self.threshold)")

    self.pre_prepare()
    self.prepare()
    self.commit()
    self.show_result()
```

```
# 1. Pre-prepare phase: All nodes (including malicious) receive correct value
def pre_prepare(self):
    print("==== Pre-prepare Phase (All nodes receive correct proposal) =====")
    for node in self.nodes:
        node.preprepare_val = self.value # Key rule: All nodes get correct value!
        status = 'Malicious' if node.is_malicious else 'Honest'
        print(f"Node {node.node_id} [{status}] received: {node.preprepare_val}")
```

```
# 2. Prepare phase: Malicious nodes broadcast fake messages
def prepare(self):
    print("==== Prepare Phase (Malicious nodes broadcast fake messages) =====")
    # Step 1: All nodes broadcast messages
    for sender in self.nodes:
        # Determine broadcast content
        if sender.is_malicious:
            # Malicious nodes broadcast fake message
            send_msg = 'fake (self.value), (sender.node_id)'
        else:
            # Honest nodes broadcast real message
            send_msg = self.value

    # Broadcast to all nodes
    for receiver in self.nodes:
        if receiver.is_malicious:
            continue # Malicious nodes don't count messages
        receiver.prepare_cnt[send_msg] += 1 # Receiver counts messages
```

```
# Step 2: Check if each honest node reaches committed state
print(f"Prepare Phase Statistics (Honest Nodes Only):")
print("Total honest nodes: (self.honest_num)")
print("Honest node ID list: (self.honest_node_ids)")
print("Prepared state check for each honest node:")
for node in self.honest_nodes:
    if node.is_malicious:
        continue
    # Count correct messages
    correct_cnt = node.prepare_cnt.get(self.value, 0)
    # Check if each threshold
    if correct_cnt == self.threshold:
        node.prepared = True
        print(f"Honest Node {node.node_id}: Correct messages {correct_cnt} > threshold {self.threshold}. Prepared state achieved!")
    else:
        node.prepared = False
        print(f"Honest Node {node.node_id}: Correct messages {correct_cnt} < threshold {self.threshold}. Prepared state not achieved!")
```

```
# 3. Commit phase: Only prepared honest nodes broadcast commit messages
def commit(self):
    print("==== Commit Phase =====")
    # Step 1: Broadcast commit messages
    for sender in self.nodes:
        if sender.is_malicious or not sender.prepared:
            continue # Malicious/prepared nodes don't broadcast
        # Broadcast correct proposal to all honest nodes
        for receiver in self.honest_nodes:
            if receiver.is_malicious:
                continue
            receiver.commit_cnt[self.value] += 1
```

```
# Step 2: Check if each honest node reaches committed state
print("Commit Phase Statistics (Honest Nodes Only):")
for node in self.honest_nodes:
    if node.is_malicious:
        continue
    commit_cnt = node.commit_cnt.get(self.value, 0)
    print(f"Node {node.node_id}: Received commit messages {commit_cnt}")
    if commit_cnt == self.threshold:
        node.committed = True
        node.result = self.value
    else:
        node.committed = False
        node.result = None
```

```
# Show final consensus result
def show_result(self):
    print("==== Final Consensus Result =====")
    honest_results = []
    for node in self.honest_nodes:
        if node.is_malicious:
            continue
        print(f"Honest Node {node.node_id} Final Result: {node.result}")
        honest_results.append(node.result)

    # Judge if consensus succeeded
    if len(set(honest_results)) == 1 and honest_results[0] is not None:
        print("✅ Consensus Success! All honest nodes reached agreement")
    else:
        print("❌ Consensus Failed! Honest nodes did not reach agreement")
```

PBFT_Simulator Class Methods:

Method

Description

commit()

Commit phase: Prepared honest nodes broadcast commit msg; count & verify commit state

```
# 3. Commit phase: Only prepared honest nodes broadcast commit messages
def commit(self):
    print("\n==== Commit Phase =====")
    # Step 1: Broadcast commit messages
    for sender in self.nodes:
        if sender.is_malicious or not sender.prepared:
            continue # Malicious/unprepared nodes don't broadcast
        # Broadcast correct proposal to all honest nodes
        for receiver in self.nodes:
            if receiver.is_malicious:
                continue
            receiver.commit_cnt[self.value] += 1
```

```
# Step 2: Check if each honest node reaches committed state
print("\n[Commit Phase Statistics (Honest Nodes Only)]:")
for node in self.nodes:
    if node.is_malicious:
        continue
    commit_cnt = node.commit_cnt.get(self.value, 0)
    print(f"Node {node.node_id}: Received commit messages={commit_cnt}")
    if commit_cnt >= self.threshold:
        node.committed = True
        node.result = self.value
    else:
        node.committed = False
        node.result = None
```

PBFT_Simulator Class Methods:

Method	Description
pre_prepare()	Pre-prepare phase: All nodes receive correct proposal value
prepare()	Prepare phase: Honest nodes broadcast real msg; malicious broadcast fake msg; count & verify prepare state
commit()	Commit phase: Prepared honest nodes broadcast commit msg; count & verify commit state
run()	Execute full PBFT consensus: run pre-prepare → prepare → commit → show result
show_result()	Display final consensus result of honest nodes; judge consensus success

```
# Run the complete PBFT consensus process
def run(self):
    print("==== PBFT Consensus Simulation Started =====")
    print("Total nodes: (self.N) | Malicious nodes: (self.M) | Honest nodes: (self.honest_num)")
    print("PBFT Security Threshold: T = (N+1)/3 = (current (self.T) + (self.N-1)/3) - (self.safe)")
    print("Consensus Threshold: (self.T) of total nodes: (self.threshold)")

    self.pre_prepare()
    self.prepare()
    self.commit()
    self.show_result()
```

```
# 1. Pre-prepare phase: All nodes (including malicious) receive correct value
def pre_prepare(self):
    print("==== Pre-prepare Phase (All nodes receive correct proposal) =====")
    for node in self.nodes:
        node.preprepare_val = self.value # Key rule: All nodes get correct value!
        status = 'Malicious' if node.is_malicious else 'Honest'
        print(f"Node {node.node_id} [{status}] received: {node.preprepare_val}")
```

```
# 2. Prepare phase: Malicious nodes broadcast fake messages
def prepare(self):
    print("==== Prepare Phase (Malicious nodes broadcast fake messages) =====")
    # Step 1: All nodes broadcast messages
    for sender in self.nodes:
        # Determine broadcast content
        if sender.is_malicious:
            # Malicious nodes broadcast fake message
            send_msg = 'fake (self.value) (sender node_id)'
        else:
            # Honest nodes broadcast real message
            send_msg = self.value

    # Broadcast to all nodes
    for receiver in self.nodes:
        if receiver.is_malicious:
            continue # Malicious nodes don't count messages
        receiver.prepare_cnt[send_msg] += 1 # Receiver counts messages
```

```
# Step 2: Check if each honest node reaches committed state
print(f"Prepare Phase Statistics (Honest Nodes Only):")
print("Total honest nodes: (self.honest_num)")
print("Honest node ID list: (self.honest_node_ids)")
print("Prepared state check for each honest node:")
for node in self.honest_nodes:
    if node.is_malicious:
        continue
    # Count correct messages
    correct_cnt = node.prepare_cnt.get(self.value, 0)
    # Check if each threshold
    if correct_cnt == self.threshold:
        node.prepared = True
        print(f"Honest Node {node.node_id}: Correct messages (correct_cnt) > threshold (self.threshold). Prepared state achieved!")
    else:
        node.prepared = False
        print(f"Honest Node {node.node_id}: Correct messages (correct_cnt) < threshold (self.threshold). Prepared state not achieved!")
```

```
# 3. Commit phase: Only prepared honest nodes broadcast commit messages
def commit(self):
    print("==== Commit Phase =====")
    # Step 1: Broadcast commit messages
    for sender in self.nodes:
        if sender.is_malicious or not sender.prepared:
            continue # Malicious/prepared nodes don't broadcast
        # Broadcast correct proposal to all honest nodes
        for receiver in self.honest_nodes:
            if receiver.is_malicious:
                continue
            receiver.commit_cnt[self.value] += 1
```

```
# Step 2: Check if each honest node reaches committed state
print("Commit Phase Statistics (Honest Nodes Only):")
for node in self.honest_nodes:
    if node.is_malicious:
        continue
    commit_cnt = node.commit_cnt.get(self.value, 0)
    print(f"Node {node.node_id}: Received commit messages (commit_cnt)")
    if commit_cnt == self.threshold:
        node.committed = True
        node.result = self.value
    else:
        node.committed = False
        node.result = None
```

```
# Show final consensus result
def show_result(self):
    print("==== Final Consensus Result =====")
    honest_results = []
    for node in self.honest_nodes:
        if node.is_malicious:
            continue
        print(f"Honest Node {node.node_id} Final Result: {node.result}")
        honest_results.append(node.result)

    # Judge if consensus succeeded
    if len(set(honest_results)) == 1 and honest_results[0] is not None:
        print("✅ Consensus Success! All honest nodes reached agreement")
    else:
        print("❌ Consensus Failed! Honest nodes did not reach agreement")
```


PBFT_Simulator Class Methods:

Method

Description

run()	Execute full PBFT consensus: run pre-prepare → prepare → commit → show result
-------	---

```
# Run the complete PBFT consensus process
def run(self):
    print("==== PBFT Consensus Simulation Started =====")
    print(f"Total nodes: {self.N} | Malicious nodes: {self.f} | Honest nodes: {self.honest_num}")
    print(f"PBFT Security Threshold:  $f \leq (N-1)/3 \rightarrow$  Current {self.f}  $\leq ((self.N-1)/3) \rightarrow$  {self.safe}")
    print(f"Consensus Threshold ( $\geq 2/3$  of total nodes): {self.threshold}")

    self.pre_prepare()
    self.prepare()
    self.commit()
    self.show_result()
```

PBFT_Simulator Class Methods:

Method	Description
pre_prepare()	Pre-prepare phase: All nodes receive correct proposal value
prepare()	Prepare phase: Honest nodes broadcast real msg; malicious broadcast fake msg; count & verify prepare state
commit()	Commit phase: Prepared honest nodes broadcast commit msg; count & verify commit state
run()	Execute full PBFT consensus: run pre-prepare → prepare → commit → show result
show_result()	Display final consensus result of honest nodes; judge consensus success

```
# Run the complete PBFT consensus process
def run(self):
    print("==== PBFT Consensus Simulation Started =====")
    self.pre_prepare()
    print("Total nodes: (self.N) | Malicious nodes: (self.f) | Honest nodes: (self.honest_num)")
    print("PBFT Security Threshold:  $f < (N-1)/3$  | Current: (self.f) < (self.N-1)/3 | (self.safe)")
    print("Consensus Threshold: (self.thresh) | (self.threshold)")

    self.pre_prepare()
    self.prepare()
    self.commit()
    self.show_result()
```

```
# 1. Pre-prepare phase: All nodes (including malicious) receive correct value
def pre_prepare(self):
    print("==== Pre-prepare Phase (All nodes receive correct proposal) =====")
    for node in self.nodes:
        node.preprepare_val = self.value # Key rule: All nodes get correct value!
        status = 'Malicious' if node.is_malicious else 'Honest'
        print(f"Node {node.node_id} [{status}] received: {node.preprepare_val}")
```

```
# 2. Prepare phase: Malicious nodes broadcast fake messages
def prepare(self):
    print("==== Prepare Phase (Malicious nodes broadcast fake messages) =====")
    # Step 1: All nodes broadcast messages
    for sender in self.nodes:
        # Determine broadcast content
        if sender.is_malicious:
            # Malicious nodes broadcast fake message
            send_msg = 'fake (self.value) (sender node_id)'
        else:
            # Honest nodes broadcast real message
            send_msg = self.value

    # Broadcast to all nodes
    for receiver in self.nodes:
        if receiver.is_malicious:
            continue # Malicious nodes don't count messages
        receiver.prepare_cnt[send_msg] += 1 # Receiver counts messages
```

```
# Step 2: Check if each honest node reaches committed state
def commit(self):
    print("==== Commit Phase Statistics (Honest Nodes Only) =====")
    # Check if each threshold
    if correct_cnt == self.threshold:
        node.prepared = True
        print(f"Node {node.node_id}: Correct messages {correct_cnt} > threshold {self.threshold}. Prepared state achieved!")
    else:
        node.prepared = False
        print(f"Node {node.node_id}: Correct messages {correct_cnt} < threshold {self.threshold}. Prepared state not achieved!")
```

```
# 3. Commit phase: Only prepared honest nodes broadcast commit messages
def commit(self):
    print("==== Commit Phase =====")
    # Step 1: Broadcast commit messages
    for sender in self.nodes:
        if sender.is_malicious or not sender.prepared:
            continue # Malicious/prepared nodes don't broadcast
        # Broadcast correct proposal to all honest nodes
        for receiver in self.nodes:
            if receiver.is_malicious:
                continue
            receiver.commit_cnt[self.value] += 1
```

```
# Step 2: Check if each honest node reaches committed state
def show_result(self):
    print("==== Final Consensus Result =====")
    honest_results = []
    for node in self.nodes:
        if node.is_malicious:
            continue
        print(f"Honest Node {node.node_id} Final Result: {node.result}")
        honest_results.append(node.result)

    # Judge if consensus succeeded
    if len(set(honest_results)) == 1 and honest_results[0] is not None:
        print("Consensus Success! All honest nodes reached agreement")
    else:
        print("Consensus Failed! Honest nodes did not reach agreement")
```

```
# Show final consensus result
def show_result(self):
    print("==== Final Consensus Result =====")
    honest_results = []
    for node in self.nodes:
        if node.is_malicious:
            continue
        print(f"Honest Node {node.node_id} Final Result: {node.result}")
        honest_results.append(node.result)

    # Judge if consensus succeeded
    if len(set(honest_results)) == 1 and honest_results[0] is not None:
        print("Consensus Success! All honest nodes reached agreement")
    else:
        print("Consensus Failed! Honest nodes did not reach agreement")
```

PBFT_Simulator Class Methods:

Method

Description

show_result()

Display final consensus result of honest nodes; judge consensus success

```
# Show final consensus result
def show_result(self):
    print("\n===== Final Consensus Result =====")
    honest_results = []
    for node in self.nodes:
        if node.is_malicious:
            continue
        print(f"Honest Node {node.node_id} Final Result: {node.result}")
        honest_results.append(node.result)

    # Judge if consensus succeeded
    if len(set(honest_results)) == 1 and honest_results[0] is not None:
        print("\n✅ Consensus Success! All honest nodes reached agreement")
    else:
        print("\n❌ Consensus Failed! Honest nodes did not reach agreement")
```


Additional Methods:

Attribute/Method Description

get_input()	Open GUI dialog to get PBFT node parameters; validate inputs and return valid total & malicious nodes
-------------	---

```
# Get user input via GUI dialog
def get_input():
    # Initialize tkinter and hide main window
    root = tk.Tk()
    root.withdraw()

    # Get total node count (must be ≥4)
    while True:
        total_input = simpledialog.askstring("Input Total Nodes", "Please enter total
number of PBFT nodes (positive integer ≥4): ", parent=root)
        if total_input is None: # User cancelled input
            exit()
        # Validate input
        try:
            total_nodes = int(total_input)
            if total_nodes >= 4:
                break
            else:
                messagebox.showerror("Input Error", "Total nodes must be ≥4!")
        except ValueError:
            messagebox.showerror("Input Error", "Please enter a valid positive integer!")
```

```
# Get malicious node count (0 ≤ malicious nodes < total nodes)
while True:
    malicious_input = simpledialog.askstring("Input Malicious Nodes", f"Total nodes=
{total_nodes}, please enter number of malicious nodes (0 ≤ number < {total_nodes}): ",
parent=root)
    if malicious_input is None: # User cancelled input
        exit()
    # Validate input
    try:
        malicious_count = int(malicious_input)
        if 0 <= malicious_count < total_nodes:
            break
        else:
            messagebox.showerror("Input Error", f"Malicious nodes must satisfy: 0 ≤
number < {total_nodes}!")
    except ValueError:
        messagebox.showerror("Input Error", "Please enter a valid integer!")

    return total_nodes, malicious_count
```

Main Programme:



```
# Main program entry
if __name__ == "__main__":
    # Get user input
    N, F = get_input()
    # Initialize simulator and run
    sim = PBFT_Simulator(N, F)
    sim.run()
```



Simulation Scenario Design



To verify the core conclusions of PBFT, two comparative scenarios are designed. Both use **9 total nodes (N=9)**.



The security threshold is $f \leq (9-1)/3 = 2$, meaning consensus succeeds if malicious nodes ≤ 2 , and fails otherwise.

Scenario	Total Nodes	Malicious Nodes	Honest Nodes	Meets SecurityThreshold	Expected Result
Scenario 1	9	1	8	Yes ($1 \leq 2$)	Consensus success; all honest nodes reach agreement
Scenario 2	9	3	6	No ($3 > 2$)	Consensus failure; no honest nodes reach final decision



Simulation Actual Results

 **Consensus Success** 

Scenario	Total Nodes	Malicious Nodes	Honest Nodes	Meets SecurityThreshold	Expected Result
Scenario 1	9	1	8	Yes ($1 \leq 2$)	Consensus success; all honest nodes reach agreement



**Pre-prepare phase**

The primary node broadcasts a correct proposal. Only 1 malicious node tampers with the proposal, while 8 honest nodes receive the correct one.

**Prepare phase**

All 8 honest nodes broadcast valid prepare messages. Each honest node receives 8 correct prepare messages, well above the threshold $(2 \times 8) // 3 + 1 = 6$, so all honest nodes confirm the proposal is valid.

**Commit phase**

All 8 honest nodes broadcast commit messages. Each honest node receives 8 correct commit messages, exceeding the threshold. Finally, all honest nodes reach consistent consensus with the original proposal value.



Simulation Actual Results

 **Consensus Failure !**

Scenario	Total Nodes	Malicious Nodes	Honest Nodes	Meets SecurityThreshold	Expected Result
Scenario 2	9	3	6	No ($3 > 2$)	Consensus failure; no honest nodes reach final decision



**Pre-prepare phase**

Some of the 3 malicious nodes tamper with the proposal. 6 honest nodes receive the correct proposal, but invalid messages from malicious nodes interfere with the statistics in the prepare phase.

**Prepare phase**

honest nodes broadcast correct prepare messages. Each honest node only receives 5 valid prepare messages, which **does not exceed the threshold** ($(2 \times 6) // 3 + 1 = 5$), so the proposal cannot be confirmed.

**Commit phase**

No honest nodes confirm the proposal, so no commit messages are broadcast. The final decision of all honest nodes is None, and consensus cannot be reached.



Core Conclusion Verification




The results of the two scenarios fully align with the core conclusion of the PBFT algorithm

the system can reach effective consensus when the number of malicious nodes $\leq (N-1)/3$, and fails to reach consensus when the number of malicious nodes $> (N-1)/3$.



This project successfully verified the conclusion through pure logical simulation and achieved the expected goals.



4 Core Challenges



Logic Cohesion

Three-phase flow integrity

Ensuring Prepare/Pre-prepare/Commit stages



Threshold Precision

2/3 calculation accuracy


Precise computation for safety thresholds



Malicious Behavior

Realistic attack simulation

Testing robustness under Byzantine faults



Broadcast Logic

Accurate message passing

Avoiding duplication and omission



Challenge 1 — Logic Cohesion


× Problem

Nodes may skip the "Prepare" phase and directly jump to "Commit", breaking consensus integrity.

+ Solution

Introduce an `accepted_proposal` gatekeeper flag to enforce the three-phase sequence.

Workflow Enforcement





Challenge 2 — Threshold Precision



The Problem

Integer division errors can inadvertently lower the safety threshold, leading to potential system vulnerabilities.



The Solution

Implement ceiling division to ensure the threshold is always rounded up, maintaining the 2/3 safety requirement.

Mathematical Definition

$$\text{threshold} = \left\lceil \frac{2N}{3} \right\rceil + 1$$

Where N = total number of honest nodes

Validation: Example Calculations

Total Nodes	2/3 Threshold	Our Formula
4	$2.67 \rightarrow 3$	$\lceil 2.67 \rceil + 1 = 3$
7	$4.67 \rightarrow 5$	$\lceil 4.67 \rceil + 1 = 5$
10	$6.67 \rightarrow 7$	$\lceil 6.67 \rceil + 1 = 7$



Result: The formula consistently rounds up to the correct safety threshold.



Challenge 3 — — Malicious Behavior



Problem: Simulate realistic Byzantine faults to test PBFT resilience



Solution: 50% chance to corrupt proposal + no Prepare/Commit votes

Honest Node Behavior



Propose: Generates valid block proposals



Prepare: Votes normally during prepare phase



Commit: Votes normally during commit phase

Malicious Node Behavior



Corrupt: 50% chance to send invalid data



Silent: No response during prepare phase



Silent: No response during commit phase

Objective: Verify consensus algorithm can tolerate up to $1/3$ malicious nodes and maintain data integrity.

Challenge 4 — Broadcast Logic



Problem

Message duplication or omission in distributed systems.



Solution


Implement **Double-loop** broadcast + **Dictionary** aggregation.



Guarantee

Accuracy and integrity of message delivery with automatic deduplication.

Star Network Broadcast Mechanism





Overview



What We Built

- ✓ Pure PBFT simulation
- ✓ 3-phase consensus process
- ✓ Verified 2/3 safety threshold




What We Learned

- Deepened theory understanding
- Python OOP programming skills
- Distributed systems thinking



Where We Go

- 5 key improvement paths
- Enhance production readiness
- Optimization for real-world use



Project Summary — What We Achieved



Core Outcome

Pure Python simulation validating PBFT's 2/3 safety threshold.
Successfully implemented all consensus phases without external dependencies.

Three Key Metrics



Code Lines:~200 (concise)



Dependencies:0 (pure logic)



Phases:3/3 fully implemented

Consensus Process Timeline



Pre-Prepare Phase



Prepare Phase



Commit Phase



Consensus Successfully Reached



Personal Learnings — 3 Dimensions




Theory

PBFT internals
Byzantine math
3-phase workflow



Practice

Python OOP
Debugging logic
Broadcast modeling



Mindset

Distributed thinking
Threshold-first design
Problem decomposition

"Continuous learning is the minimum requirement for success in any field."



Future Work — 5 Directions

Future Work Directions

01. View Change

Impact: Dynamic leader election mechanism

02. Rich Malice

Impact: Simulate equivocation & selective silence

03. Visualization

Impact: Real-time message flow graphs

04. Multi-Proposal

Impact: Concurrent request handling capability

05. Optimization

Impact: Scale system to support 100+ nodes

Development Roadmap

Phase 1: View Change

Implementing dynamic leader election.

Phase 2: Rich Malice

Adding advanced malicious node behaviors.

Phase 3: Visualization

Building real-time message flow graphs.

Phase 4: Multi-Proposal

Enabling concurrent request handling.

Phase 5: Optimization

Scaling to support 100+ nodes.



References & Closing

Key Reference

Castro, M., & Liskov, B. (1999). *Practical Byzantine Fault Tolerance*. Proceedings of OSDI.

Additional References

- Blockchain Technology Principles and Applications (2nd Edition). Cai Weide et al.
- Python Object-Oriented Programming. Dong Fuguo.

Closing Statement

“A clean, verifiable simulation that bridges theory and practice — foundation for deeper distributed systems exploration.”



Introduction


Theory
Foundations

Code

Results

Challenge

Summary



Thanks For **L**istening!